

BIGDATA 210: Introduction to Data Engineering

Autumn 2018

Module 5: Programming Spark Part III

Jerry Kuch

jkuch@uw.edu

Week 5 Agenda

- VM/Docker/Sandbox Login Issues?
- Last Week: *Programming in Spark Part II*:
 - Partitioning and Shuffling
 - Persistence and Caching
 - Spark SQL, DataFrames, Datasets, Hive integration
 - We'll finish the very last bit off tonight...
- *Programming in Spark Part III*
 - Shared Variables
 - Spark UI
 - Spark SQL Server
 - Notebooks (Jupyter and Zeppelin)
 - Summary

First off: Any Docker/VM Issues

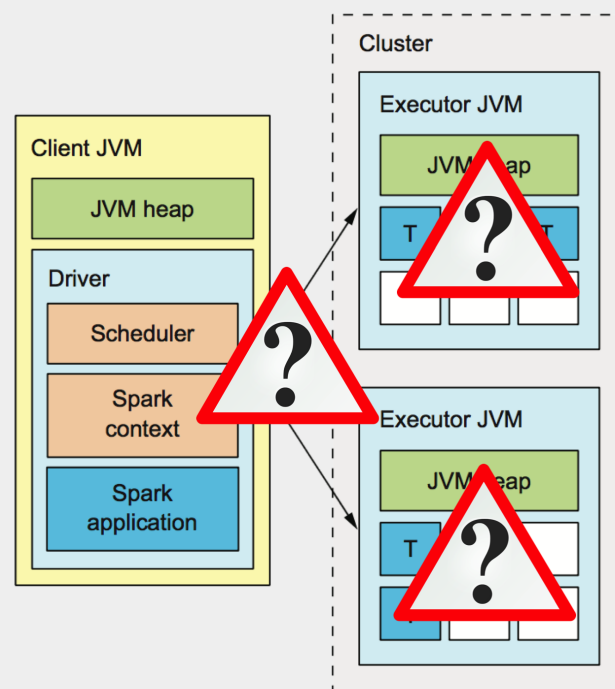
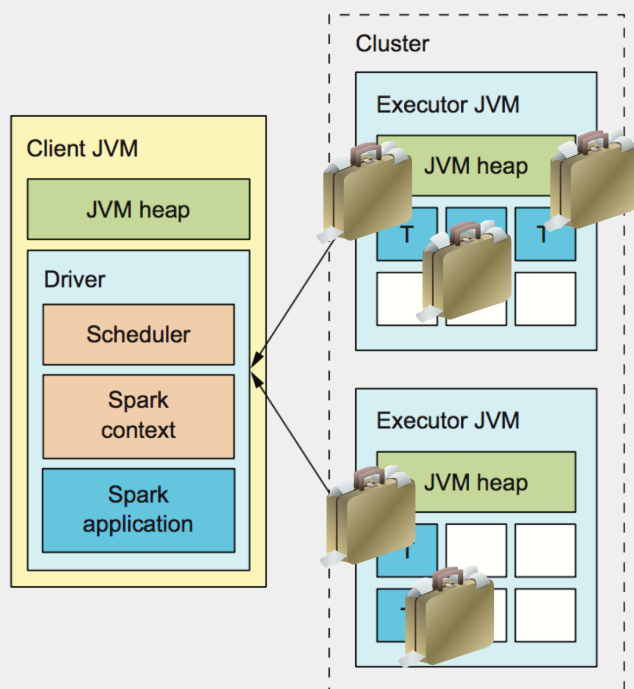
- Any Issues?



Spark Shared Variables


Broadcast Variables
and
Accumulators

Spark Shared Variables



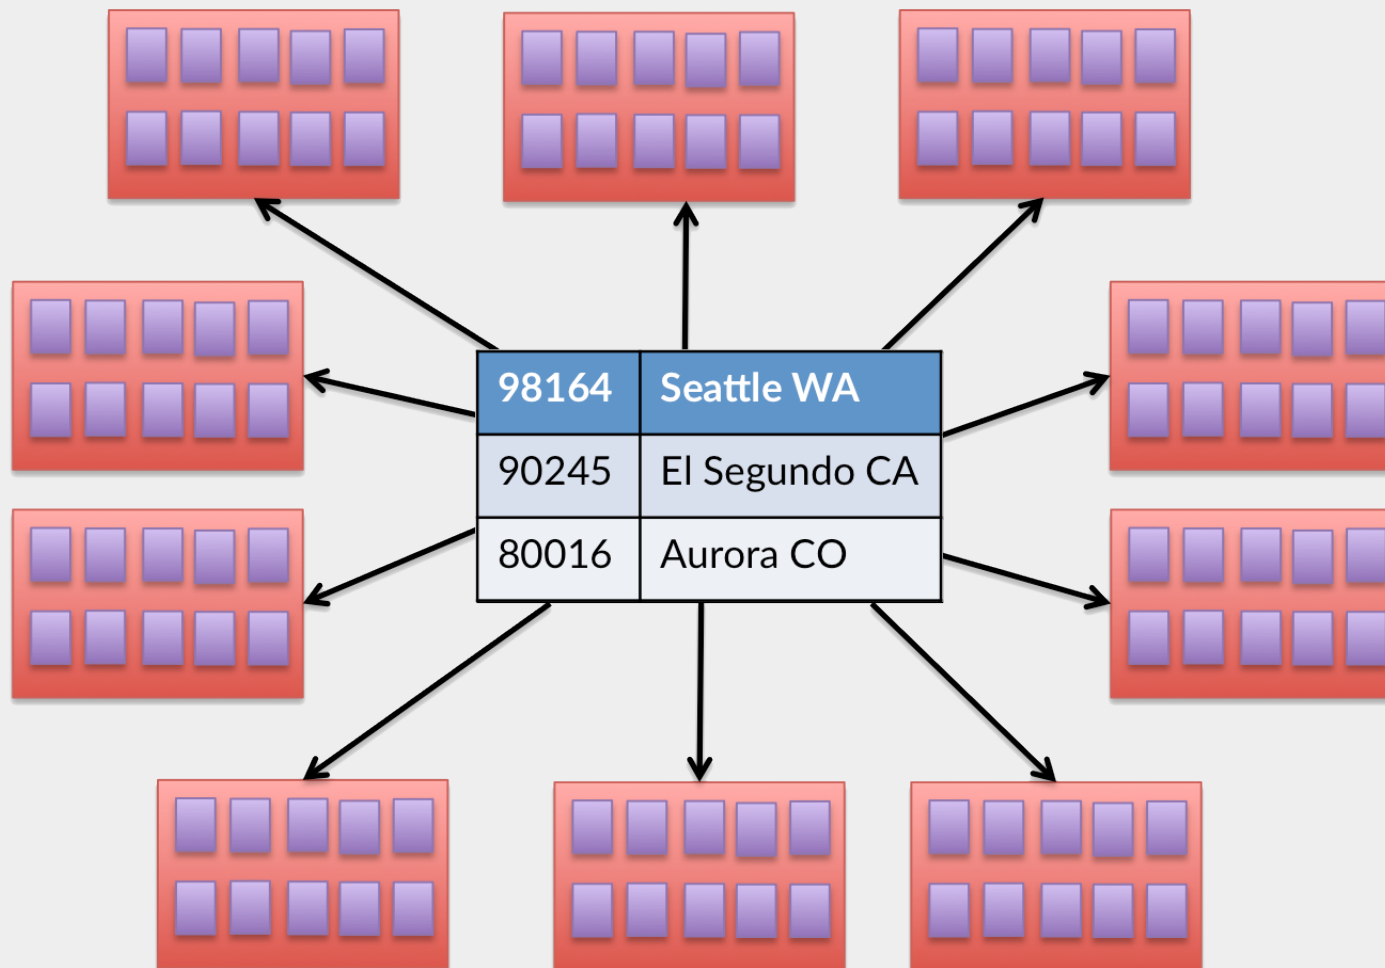
Broadcast Variables

- Cache read-only variable on each node



e.g. A Map stored as regular Scala Map sent to all tasks individually

Broadcast Variables



Map stored as broadcast variable and sent only once to each node

Broadcast Variables

```
val broadcastMap = Map("98164" -> "Seattle WA", "90245" -> "El Segundo CA", "80016" -> "Aurora CO")
```

```
val broadcastVar = sc.broadcast(broadcastMap)
```

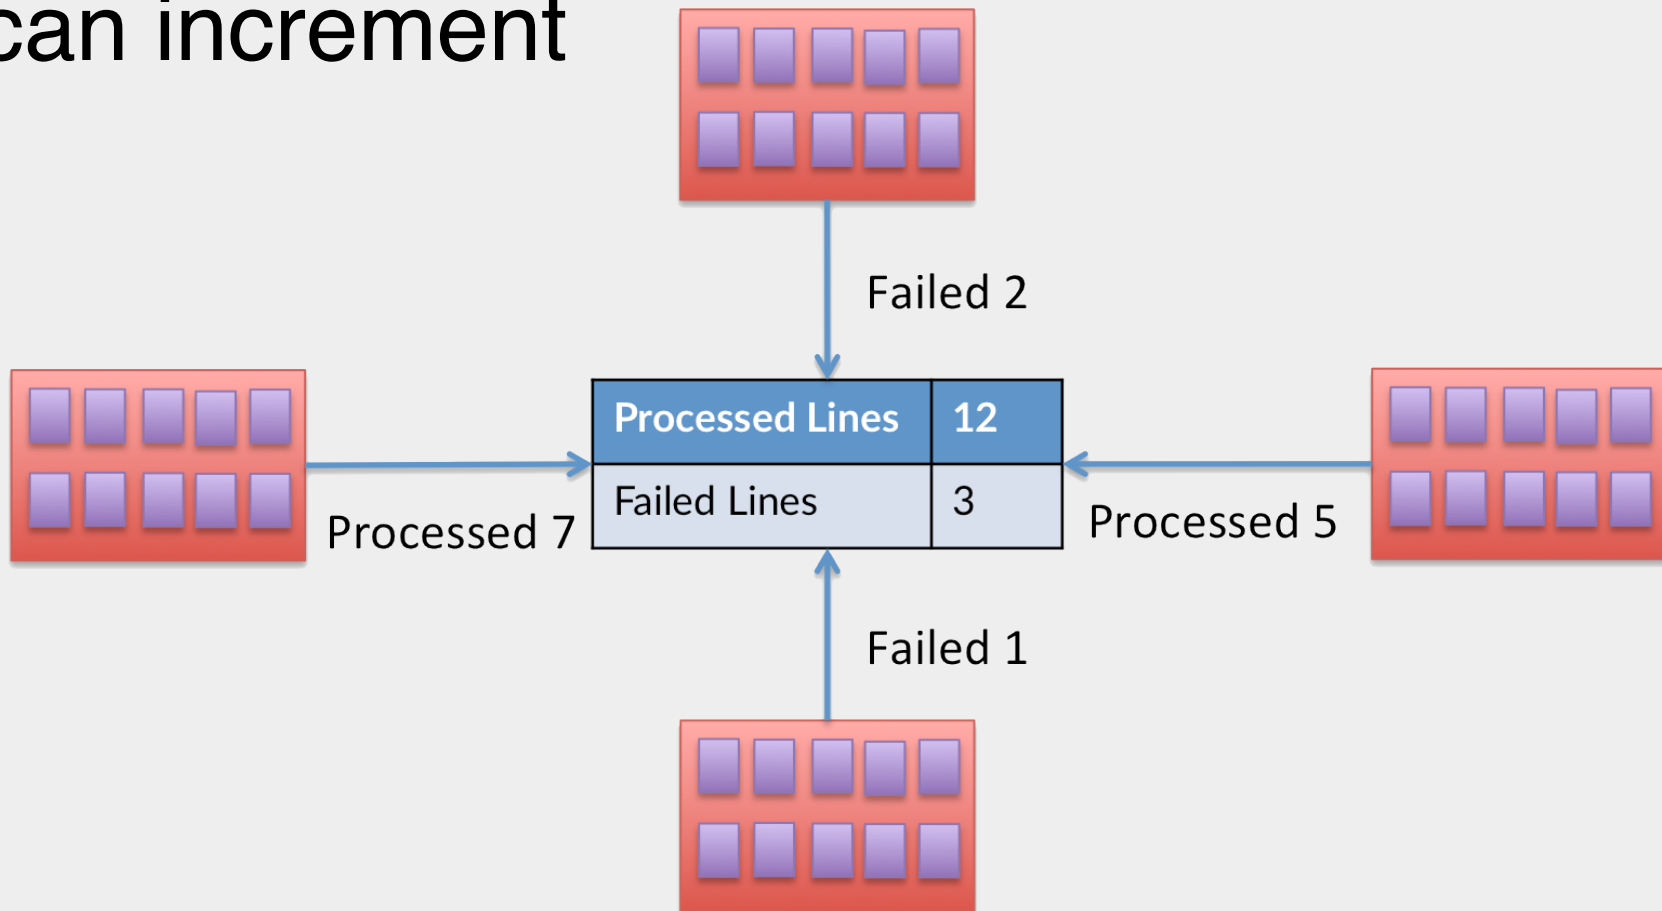
```
val data = sc.parallelize(Array("98164", "90245", "80016"))  
val transformed = data.map(v => (broadcastVar.value(v)))  
transformed.collect()
```

```
Array[String] = Array(Seattle WA, El Segundo CA, Aurora CO)
```

```
broadcastVar.unpersist
```


Accumulators

- Shared variable *on the driver* that all tasks can increment



Four executors writing to accumulator stored in the driver

Accumulators

```
val processed = sc.longAccumulator("Processed Lines")  
val failed = sc.longAccumulator("Failed Lines")
```


```
val rdd = sc.parallelize(1 to 12)  
rdd.foreach(x => processed.add(1))
```

```
processed.value
```

```
Long = 12
```

Spark Web UI

Spark Web UI

 2.1.0

Jobs | Stages | Storage | Environment | Executors | SQL

IBM Spark Kernel application UI

Spark Jobs (?)

User: ubuntu
Total Uptime: 26.7 h
Scheduling Mode: FIFO
Completed Jobs: 30

[Event Timeline](#)

Completed Jobs (30)

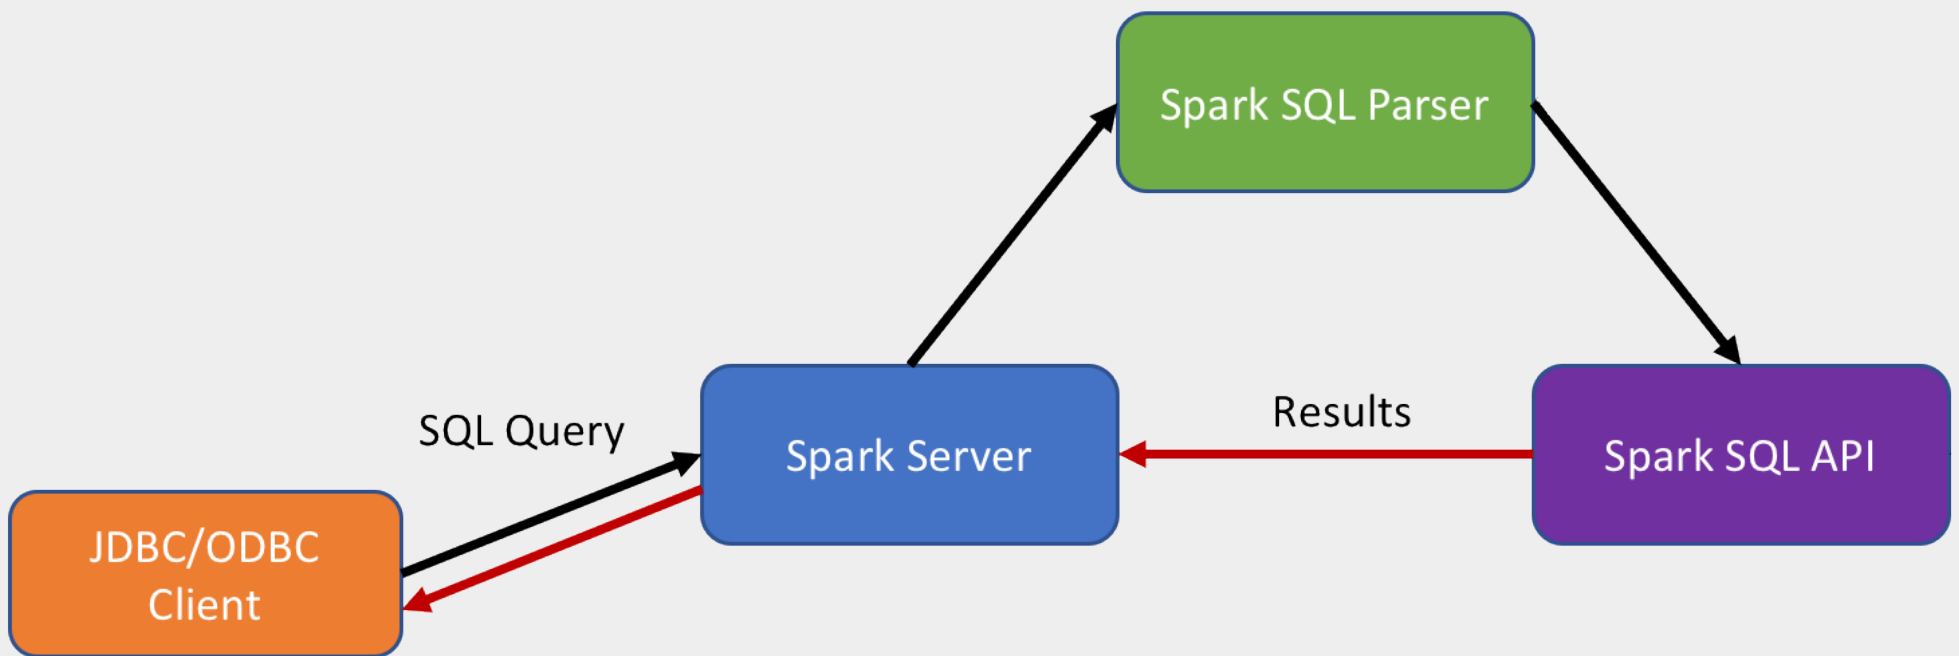
Job Id ▾	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
29	count at <console>:53	2017/04/02 22:14:08	0.6 s	2/2 (5 skipped)	201/201 (412 skipped)
28	show at <console>:51	2017/04/02 22:14:07	0.4 s	1/1 (5 skipped)	200/200 (412 skipped)
27	approxQuantile at QuantileDiscretizer.scala:151	2017/04/02 22:14:07	0.4 s	1/1 (5 skipped)	200/200 (412 skipped)
26	approxQuantile at QuantileDiscretizer.scala:151	2017/04/02 22:14:06	0.4 s	1/1 (5 skipped)	200/200 (412 skipped)
25	approxQuantile at QuantileDiscretizer.scala:151	2017/04/02 22:14:05	0.3 s	1/1 (5 skipped)	200/200 (412 skipped)
24	approxQuantile at <console>:42	2017/04/02 22:14:05	0.3 s	1/1 (5 skipped)	200/200 (412 skipped)
23	approxQuantile at <console>:42	2017/04/02 22:14:04	0.3 s	1/1 (5 skipped)	200/200 (412 skipped)
22	approxQuantile at <console>:42	2017/04/02 22:14:03	0.4 s	1/1 (5 skipped)	200/200 (412 skipped)
21	describe at <console>:42	2017/04/02 22:14:02	0.5 s	2/2 (5 skipped)	201/201 (412 skipped)
20	count at <console>:42	2017/04/02 22:14:01	0.5 s	3/3	205/205
19	run at ThreadPoolExecutor.java:1142	2017/04/02 22:14:01	0.3 s	1/1 (5 skipped)	200/200 (412 skipped)
18	count at <console>:32	2017/04/02 22:14:00	0.5 s	3/3	205/205
17	count at <console>:38	2017/04/02 22:14:00	0.4 s	2/2 (5 skipped)	201/201 (412 skipped)
16	show at <console>:42	2017/04/02 22:13:59	8 ms	1/1 (1 skipped)	3/3 (4 skipped)
15	show at <console>:42	2017/04/02 22:13:59	11 ms	1/1 (1 skipped)	4/4 (4 skipped)

http://driver.ip:4040-?

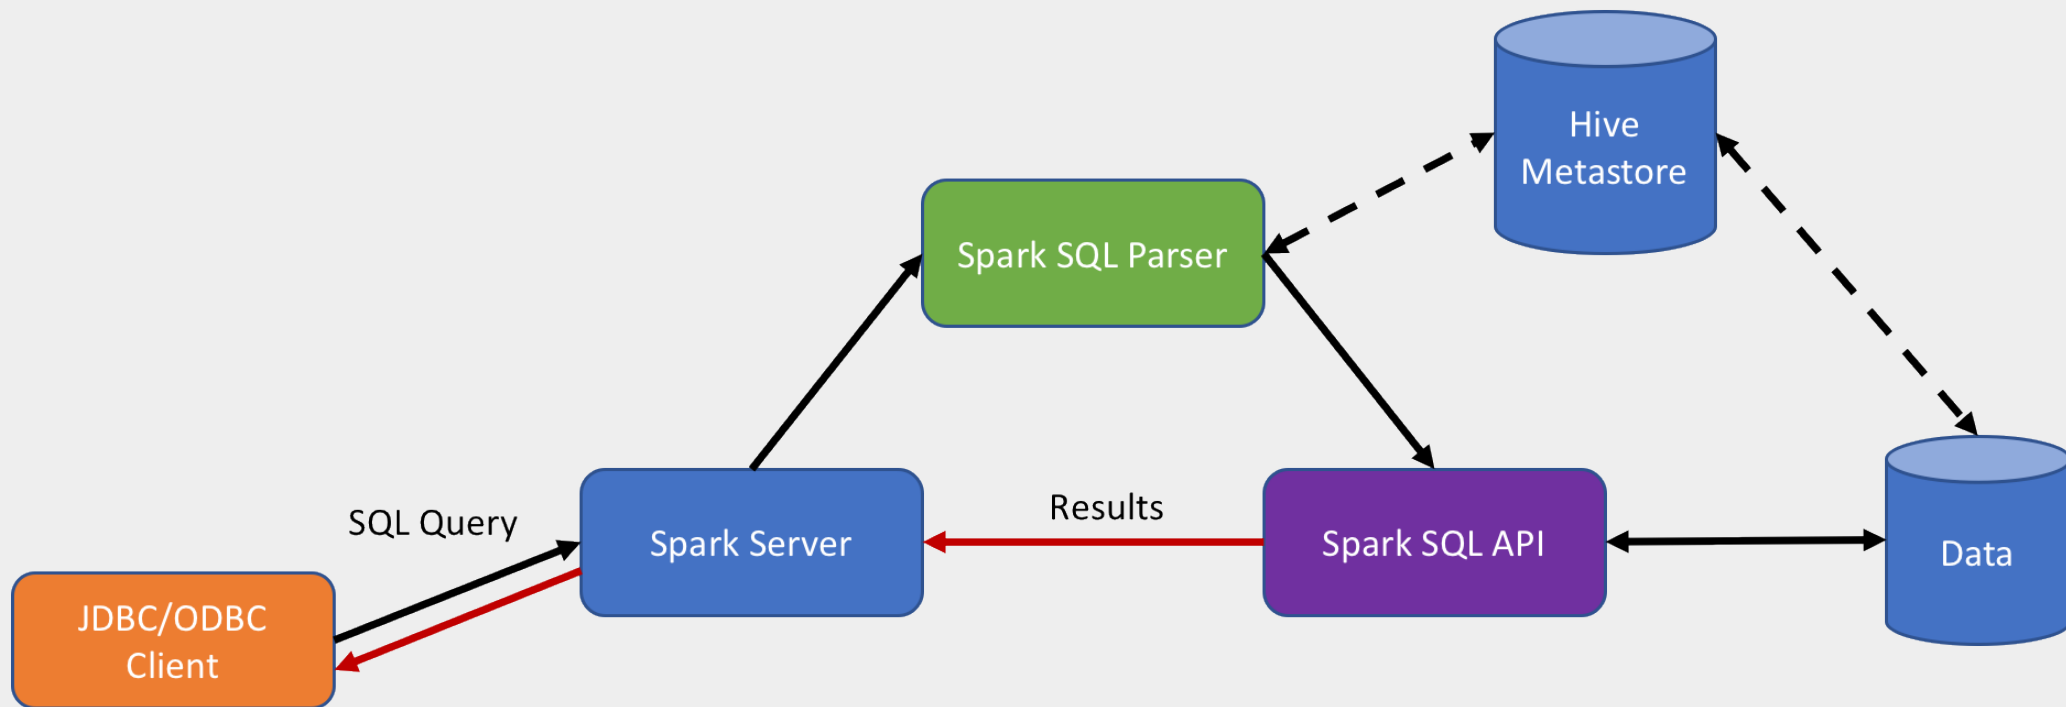
Spark SQL

Thrift Server, Hive
and Command Line

Spark SQL Thrift Server



Spark SQL Thrift Server: With the Hive megastore...



Spark SQL Command Line

- Run 'spark-sql' in the Sandbox
 - Connects to Hive metastore / Thrift Server
 - Starts up Web UI similar to spark-shell on indicated port ~4040 or so
- Then do SQL queries!
- Observe what happens in the UI

Notebooks

Jupyter and Zeppelin

Notebooks: Jupyter and Zeppelin

- *Notebooks* are another front end to various runtime environments
 - Interact with user via web browser
 - Backend environments provided by *kernels* that expose various runtimes (e.g. Scala, Spark, Python)
- Our universe contains two major players:
 - Jupyter
 - Zeppelin

Jupyter Notebook

- Python based notebook platform
- Formerly “IPython”
- Supports lots of “*kernels*” besides Spark
 - Spark Scala support can be buggy
 - Needs special installation
- To use PySpark, see week 5 assignment for setup and usage

Zeppelin Notebook

- Part of “Hadoop ecosystem” (Apache)
- Already set up in your sandbox
- Accessible through Ambari UI
- Lots of “big data” tools
 - Flink, Cassandra, Beam

Notebooks: Pros and Cons

- **Pros:**
 - Convenient
 - Easy scratch pad or workspace
 - “Self documenting”
- **Cons:**
 - Statefulness can yield confusion, especially w.r.t. intermediate results
 - Editing code is clunky and primitive (compared to good IDEs)
 - Some kernels are buggy

The “Just Enough Scala for Spark” Notebook

- A Jupyter notebook
- Available in a Docker container a few weeks back...

Module 5: Summary

Programming in Spark, Part III

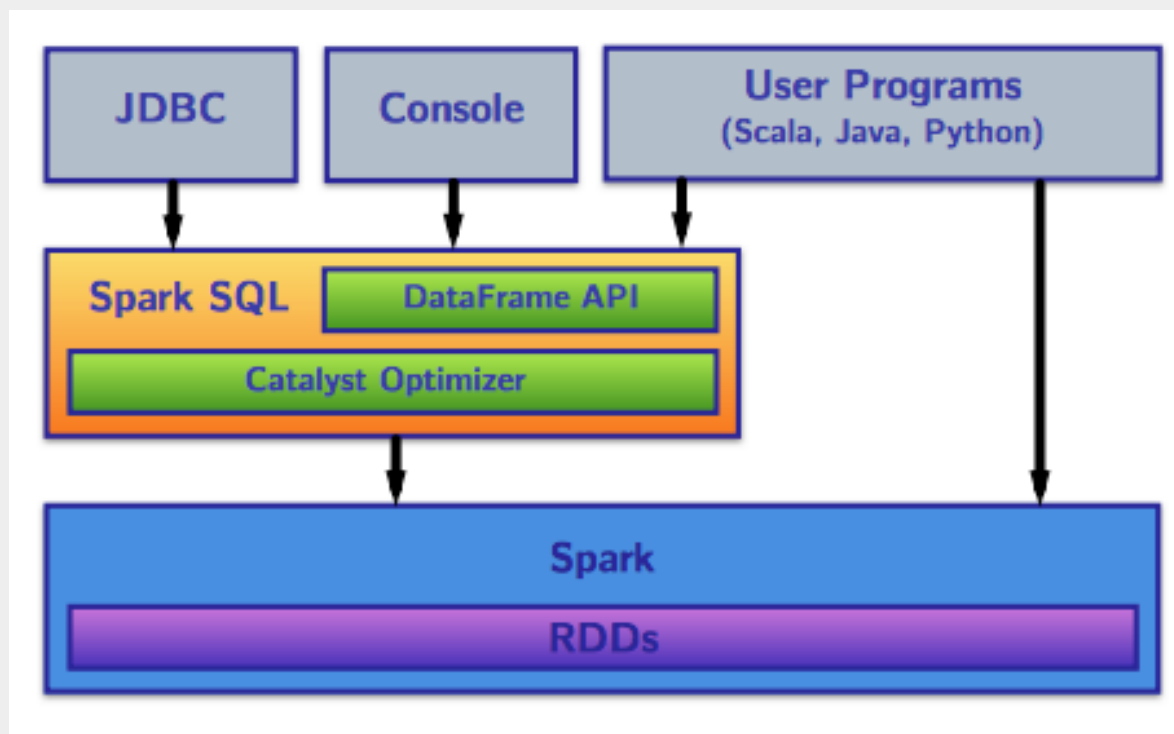
Summary (this week and last)

- Spark SQL provides three main things:
 - SQL literal syntax
 - Allows us to mix SQL queries with Scala/Spark
 - Possibility of database style optimizations
 - DataFrames
 - Spark SQL's core abstraction
 - Basically an RDD full of records with a known schema—schema required, unlike with RDDs
 - **Untyped**, i.e. Scala compiler doesn't check types in the schema! You can hit runtime exceptions!
 - Register a DataFrame as a temporary SQL view to query
 - DataSets
 - Provides Spark SQL options plus **type safety**
 - In reality: **type** DataFrame = Dataset[Row]
 - Are **typed** distributed collections of data
 - Require structured/semi-structured data, like DataFrames, but can mix and match DataFrame and RDD APIs
 - DataSets kind of lie “in the middle” between DataFrames and RDDs
 - Can use relational DF operations on them
 - Datasets add more typed operations; and you can use map, flatMap, filter

Programming in Spark, Part III

Summary (this week and last)

- Broadcast Variables and Accumulators
- Notebooks: Jupyter and Zeppelin
- Spark SQL's pieces:



Programming in Spark, Part III

Summary (this week and last)

- See the references in this week's resources, especially <https://spark.apache.org/docs/latest/sql-programming-guide.html>
- *Spark in Action*, chapter 5 has good coverage of Spark SQL
- Spark SQL in its modern form, DataFrames and Datasets are fairly new