

# Introduction to Data Engineering

---

## Week 9: Distribution is Challenging

# Today's Lecture...

---

- > Distribution of data and compute:
  - > Why do it?
  - > How did it affect the designs of things we've studied?
- > Review of some Key Distribution Challenges:
  - > Replication
  - > Partitioning
  - > Transactions
  - > Consensus



# Distribution: Why and What?



# Distribution: Why Do It?



*Why do we distribute data?*

> **Scalability:**

- > *Go Bigger:* Data too big to fit on a single machine.
- > *Go Faster:* Desired read/write rates too fast for a single machine.

> **Fault Tolerance / High Availability:**

- > Keep working if machine(s) go down via redundancy

> **Latency:**

- > Keep copies of data physically nearer to its users
- > Cache results of costly operations between user and source

# Distribution: Sharing



## *Types of Sharing*

### > ***Shared-Memory on one Node:***

- > Buy a big machine; Fault tolerance may be limited and costly

### > ***Shared Distributed-Memory:***

- > Data spread over memories on multiple nodes

### > ***Shared-Disk:***

- > Array of disks connected over fast network
- > Scalability limited by contention, locking

### > ***Share Nothing:***

- > Scale horizontally over independent nodes with own CPU/RAM/disk
- > Coordinate at software layer

# Distribution: Replication and Partitioning



## *How Data Gets Distributed Over Nodes*

### > **Replication:**

- > Keep a copy of same data on multiple nodes
- > Redundancy guards against node failure/unavailability
- > Might also improve performance

### > **Partitioning:**

- > Break large data into smaller partitions
- > Assign partitions to separate nodes (sharding)

# Replication



# Replication: Challenges and Opportunities

---

## *What and Why?*

- > **Replication:** keep copies of the same data on multiple nodes, connected via a network.
- > **Why?**
  - > Keep data near the user to *reduce latency*
  - > Keep system working even with failed nodes to *increase availability*
  - > To increase number of machines that can serve reads to *increase read throughput*



# Replication: How Immutability Helps



*Why can Immutability help?*

- > **Immutability:** if replicated data doesn't change over time it is easier to:
  - > Maintain consistency
  - > Support concurrent access
- > Almost everything hard about replication involves handling **changes** to the replicated data.
- > There are several approaches to replicating data between nodes.

# Replication: Leaders and Followers



## *One Approach: Leaders and Followers*

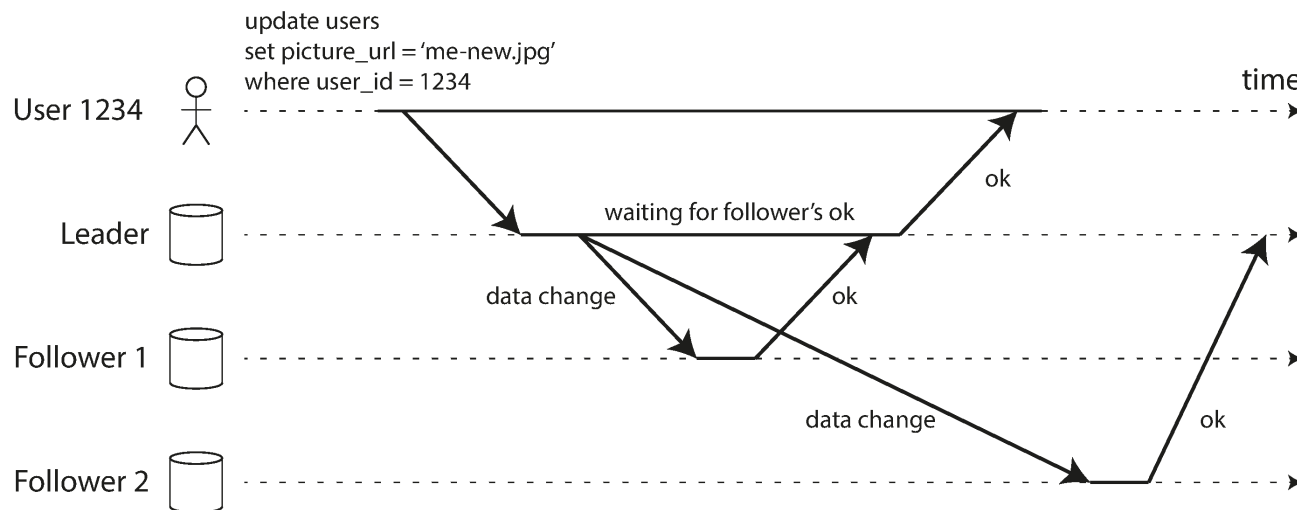
- > Each node that stores a copy of a datum is a **replica**
- > Every write should eventually reach all desired replicas.
- > Common approach: **Leader and Follower**
  - > One replica designated a leader (or master or primary)
  - > Clients submit *writes* to the leader, which writes locally first
  - > Leader sends changes to the followers, which then update their local storage, applying writes *in same order* as at leader.
- > Sometimes have a long term designated leader; other times not.

# Replication: Synchrony and Asynchrony

## Leaders and Followers: Synchrony and Asynchrony

### > Terminology: Synchronous vs. Asynchronous Updates

- > **Synchronous:** Don't report success until replicas updated.
- > **Asynchronous:** May report success before some or all replicas are written.
- > **Semi-synchronous:** some followers synchronous, others not.



# Replication: Handling Failed Replicas



## *Leaders and Followers: Handling Failed Replicas*

### > ***Handling Failed Replicas***

- > What if a leader fails? How do we tell and what do we do?
  - > Become unavailable?
  - > Decide on a new leader?
  - > What about stuff that was only on leader at failure time?
- > What if new followers are added?
  - > How do they catch up?
- > Different systems do different things, sometimes it's configurable.

# Replication: Consistency



## *Leaders and Followers: Consistency*

- > We saw earlier that replication could lag at some nodes.
- > Result: consistency issues.
- > **Consistency Levels**
  - > **Read-After-Write:** i.e. “read your writes”
  - > **Monotonic Reads:** you can’t see reads backward in time.
  - > **Consistent Prefix Reads:** if sequence of writes happens in a certain order, any reader sees them in the same order.

# Replication: Multiple Leaders



## *Multi-Leader Replication*

- > Single leader systems have downsides:
  - > Bottleneck for write throughput
  - > Single point of failure
- > Multi-leader lets writes be accepted at multiple nodes
- > Replication must still happen from leaders to followers
- > Multiple reasons to do it
- > Schemes can be complex
- > Other names: master-master, active/active replication

# Replication: Multiple Leaders



## *Multi-Leader Replication: Use Cases*

- > Multi-datacenter operation
  - > Replicas in multiple datacenters
  - > Leader(s) in each datacenter
- > Challenges:
  - > Between datacenters, latency may be high
  - > Failure may force datacenters to operate independently for a while
  - > Network conditions likely worse between DCs than within
  - > Handling write conflicts: many approaches with different pros-cons

# Replication: Leaderless Replication



## *Leaderless Replication*

- > Approach 1: Client sends write to several replicas.
- > Approach 2: A coordinator node replicates on behalf of the client
  - > Coordinator doesn't enforce a particular ordering of writes
  - > Coordinator may read from multiple replicas
  - > Related ideas: *read-repair*, *hinted-handoff*, *anti-entropy*.
  - > Quorum reading and writing: can trade off latency for consistency



# Partitioning



# Partitioning



## *Partitioning:*

- > Break data into multiple chunks spread across nodes
- > Partitions sometimes called shards, regions, tablets, vnodes, vBuckets...
- > Once data is partitioned:
  - > It can be bigger than any single node could hold
  - > Computation can be independently executed on partitions
- > Partitioning is usually combined with replication

# Partitioning: Handling Key-Value Data



## *Partitioning Key-Value Data*

### > **Goal:**

- > Spread out the data distribution and query load
- > Avoid imbalanced partitions and hot spots

### > **Schemes:**

- > Partition by key range (bad if keys very non-uniform)
- > Partition by hashes of keys (good hash function required) then give each partition a range of hashes
  - > Lose efficient range queries unless we do extra work

# Partitioning: Handling Key-Value Data



*Hot Spots can Still Occur Naturally...*

- > **e.g.** Imagine a social media site where there are celebrities and schmoes:
- > A celebrity may form a natural hot spot
- > Can use tricks, e.g.
  - > Figure out how many pieces a celebrity should be broken into...
  - > Augment their key with a piece number...
  - > ...then hash that.

# Partitioning: Handling Key-Value Data



## *Rebalancing Partitions*

- > Over a DB's lifetime things may change:
  - > Query throughput increases so we want more nodes
  - > Dataset size increases so we want more nodes
  - > Nodes fail so we want failover or replacement
- > All the above require data and queries to be *rebalanced*
- > Desiderata:
  - > After rebalancing things should be shared fairly
  - > While rebalancing system should remain available
  - > We should move as little data as possible

# Partitioning



## *Rebalancing Partitions: Strategies*

- > **Bad Idea:**  $\text{hash}(\text{data}) \bmod N$ 
  - > We change  $N$  and everything has to move
- > **Fixed Number of Partitions**
  - > Fix a number of partitions,  $N$ , where  $N \gg |\text{nodes}|$
  - > If a node is added it can steal a few partitions from every existing node; leaving nodes can donate back.
  - > Only entire partitions move between nodes
  - > Can leave old assignment of partition in place while data is moving

# Partitioning



## *Rebalancing Partitions: Strategies*

### > **Dynamic Partitioning**

- > Create partitions dynamically
- > When a partition gets too big, split it in halves
- > When a partition shrinks below threshold, merge with adjacent partitions
- > Each partition is assigned to a node
  - > Merges can happen locally
  - > Splits can overflow to other nodes to rebalance

# Partitioning: Routing Requests



## *Request Routing*

- > When a client requests data how does it know which node matters?
- > This is a *service discovery* problem with various approaches:
  - > 1. Let clients contact any node, forward requests and replies if they contacted the wrong node
  - > 2. Use a routing tier to direct clients to correct node
  - > 3. Force clients to be aware of where things are



# Partitioning



## *Request Routing*

- > How do we get queries to the correct destination node?
- > All participants must agree on:
  - > What is where?
  - > What is alive?
- > Many systems use a separate **coordination service** for this, e.g. Zookeeper in Kafka and HBase
- > Cassandra and its relatives use a **gossip protocol** among nodes to disseminate changes in cluster state

# Transactions



# Transactions



## *What and Why?*

- > Historically, transactions simplify the programming model in the face of:
  - > Failures and crashes
    - > Updates may not get everywhere
    - > Updates may be incompletely applied at a node
  - > Concurrent accesses
    - > Readers may observe partial changes in progress
    - > Writes may collide and produce corrupt results
    - > Check and act operations may race
- > Common argument: “every programmer understands transactions, so let’s just solve all our problems with them.”
  - > No.

# Transactions: Idea and Desiderata



## *Transaction: Notion and Properties*

- > A **transaction** lets an application:
  - > Group multiple reads and writes...
  - > ...in a single logical unit...
  - > ...which either commits, or aborts.
- > Traditionally transactions were **ACID**:
  - > **Atomic**: all or nothing effects
  - > **Consistent**: never leave database in a corrupt state
  - > **Isolated**: concurrent TXNs *appear* to run serially
  - > **Durable**: a committed TXN's result is stored safely

# Transactions: Scope of a Transaction



## *Single vs Multiple Object Operations*

### > **Single Objects:**

- > Atomicity and isolation are easier
- > e.g. one database row in an RDBMS; one JSON doc in a doc DB

### > **Multiple Objects:**

- > Harder
  - > May have to use *distributed locks*
  - > Failure more likely and pernicious, latency worse
  - > Many distributed datastore abandon multi-object transactions entirely
- > Scary Campfire Story: X/Open XA spec. for distributed transaction processing

# Transactions: Isolation Levels

## *Isolation Levels*

- > TXNs touching different data can safely run in parallel.
- > But, we may have problems, if two TXNs:
  - > Modify the same data, or...
  - > One reads data the other is modifying.
- > The 'I' in ACID is about hiding such problems.
- > There are different levels of isolation and:
  - > They are hard to understand
  - > Implemented inconsistently between databases.

# Transactions: Serializability



## *Serializable Isolation: The Ideal*

- > The end result of transactions, even executing in parallel, is the same as if they had executed, one at a time, *serially*, without any concurrency.
- > Thus:
  - > If the transactions behave correctly run individually...
  - > ...they continue to be correct if run concurrently.
  - > The database prevents all possible race conditions.
- > So why not demand serializability always and everywhere?

# Transactions: Serializability



## *Serializable Isolation: The Reality*

- > **Supporting Serializability is:**
  - > Expensive
  - > Complex, especially with multi-node systems
- > **Common Techniques to support Serializability:**
  - > Literally execute transactions in serial order
  - > *Two Phase Locking* (only option for many years)
  - > Optimistic Concurrency Control



# Transactions



## *Other Isolation Levels*

- > If we don't want to pay for Serializability what can we do?
- > There are ***weaker isolation levels***, but:
  - > Each doesn't guard against certain concurrency anomalies
  - > Your application will have to deal with the remainder, e.g. using explicit locking
- > **Isolation levels to read up on:**
  - > Read Committed
  - > Snapshot Isolation / Repeatable Read

# Consensus



# Consensus



## *The Consensus Problem*

- > **Consensus:** getting all nodes to agree on something.
- > Consensus is hard to do when there are:
  - > Network failures
  - > Process failures
- > Many problems are *reducible* to consensus:
  - > Leader election
  - > Atomic commit
  - > Total order broadcast
  - > Distributed locking
  - > Uniqueness constraint enforcement
  - > Membership/coordination service
- > Consensus mechanisms are complex and hard to implement

# Consensus




## *Consensus in Practical Systems*

- > We don't have time to say much about Consensus here
- > Several Consensus algorithms exist:
  - > e.g. Paxos, Raft, Viewstamped Replication
- > There are systems that implement variations of it:
  - > As a service usable by other systems
  - > e.g. Zookeeper, Raft, Chubby, etcd, consul.
  - > That are used by some systems seen in this program (e.g. HDFS, HBase, Kafka...)
- > There's debate over whether it should be a *service* or a *library*


# Consensus

## *Zookeeper: An Example Consensus System*

- > Uses a protocol called Zookeeper Atomic Broadcast
- > Deployed as a cluster of nodes (available when majority up):



- > Exposes hierarchical namespace:



# Consensus



## *Zookeeper: Exposed Abstractions*

- > **ZNodes:** each can have associated...
  - > Children
  - > Small amounts of data: e.g. status, configuration, location
  - > Each node is read and written atomically
- > **Ephemeral znodes:** exist as long as client session that created them remains active.
- > **Sequence nodes:** monotonically increasing number appended to requested name.
- > Nodes can have **watches**, that are triggered on node changes, with client notified.

# Consensus



## *Zookeeper: Guarantees*

- > **Sequential Consistency:** updates from a client will be applied in the order they were sent.
- > **Atomicity:** Updates are all or nothing.
- > **Single System Image:** client sees same view regardless of node it connects to.
- > **Reliability:** Updates survive until overwritten.
- > **Timeliness:** Client view of system guaranteed to be up to date within a certain time bound.

# Summary:

---

- > Systems are distributed because we want things like:
  - > Scalability, fault tolerance, low latency
- > Distribution requires a system to grapple with:
  - > Replication
  - > Partitioning
  - > Transactions
  - > Consensus
- > Systems we studied:
  - > Take various approaches to dealing with the above challenges...
  - > ...yielding various quirks and compromises
- > Understanding why these things are the way they are makes them easier to deal with.





# For Next Week: Project Presentations

---

- > *Does there exist anybody who has not yet:*
  - > *Formed a team?*
  - > *Chosen a dataset?*
  - > *Got started?*
- > *We'll do presentations in class next week*
  - > *If we can't fit them all in, we may have to overflow to do some by Zoom at another time*
- > *Project work can continue for another week or two (until grades due) after next week*



# References

---

- > *Designing Data-Intensive Applications*, Kleppmann, Chapters 5—9
- > <http://zookeeper.apache.org>



# Q&A



# W

# UNIVERSITY OF WASHINGTON

---

## Studying at UW

- > *What defines the students and faculty of the University of Washington? Above all, it's our belief in possibility and our unshakable optimism. It's a connection to others, both near and far. It's a hunger that pushes us to tackle challenges and pursue progress. It's the conviction that together we can create a world of good. And it's our determination to Be Boundless. Join the journey at [uw.edu](http://uw.edu).*